

Of queues and cures: A solution to modelling the inter time arrivals of cloud outage events

Jonathan Dunne^{*1} and David Malone¹

¹Maynooth University (Hamilton Institute, Maynooth University, Ireland)

^{*}Email: jonathan.dunne.2015@mumail.ie, David.Malone@nuim.ie

Abstract: The management of Cloud based outages represents a challenge for Small Medium Enterprises (SMEs), due to the variety of ways in which production outages can occur. We consider the inter-arrival times for outages events in a framework where these arrival times are used to align Systems Operations resources. Using an enterprise dataset, we address the question of how inter-arrival times are distributed by testing against a number of common distribution types. The proposed framework can help SMEs to manage their limited resource workflows. We shall also consider correlation between arrival times.

Keywords: Distribution fitting, goodness of fit, correlation resource planning.

Introduction

For the European SME the adoption of cloud technology no easy task. Due to resource constraints and myriad of failure patterns, SMEs face challenges in providing a reliable and stable service platform for their customer's needs. In this paper we describe a framework, that the SME can leverage to best manage their limited set of resources as part of incoming outage events in their cloud infrastructure.

Data Set

The study presented in this paper examines approximately 250 cloud outage events from a large enterprise system. Our study aims to answer a key question: Which distribution is best suited to model the interarrival time of cloud outage events. To answer this question a number of common distribution types were modelled; lognormal, gamma, Weibull, exponential, logistic, loglogistic and Pareto.

Results

Each distribution was modelled against the interarrival times of the dataset using the R computer package and the fitdistrplus library which also calculates the distribution parameters. Using a second package ADGofTest, the parameters of each distribution validated for their goodness of fit. Table 1 summarises the results of this test.

Table 1: Summary of Anderson-Darling GoF statistics.

Distribution Name	AD statistic	p-value
lognormal	3.039	0.026
gamma	6.034	9.347e-04
Weibull	0.975	0.371
exponential	3.110	0.024
logistic	12.819	2.765e-06
loglogistic	1.823	0.115
Pareto	0.661	0.592




Figure 1: Four goodness-of-fit plots for Weibull, loglogistic and Pareto distributions fitted to the interarrival times from the cloud outage data set.

Discussion

Table 1 shows that Pareto has the best p-value for the Anderson-Darling test, followed by loglogistic and Weibull. All other distributions were rejected as part of hypothesis testing. Figure 1 shows goodness of fit plots for the three best fitted distributions. The quantile-quantile plot shows that the Pareto distribution more closely models the data set even for extreme values.

Conclusion

It was found that the Pareto distribution is a useful distribution for modelling the interarrival times of cloud outages. This result can be used by SMEs as an arrival time parameter for a queuing model for cloud outage events.

References

Delignette-Muller, M.L. and Dutang, C. and Pouillot, R. and Denis, J.B. (2015).

Web Page <https://cran.r-project.org/web/packages/fitdistrplus/index.html>

Bellosta, C.J.G (2011). <https://cran.r-project.org/web/packages/ADGofTest/index.html>